

Regression Model

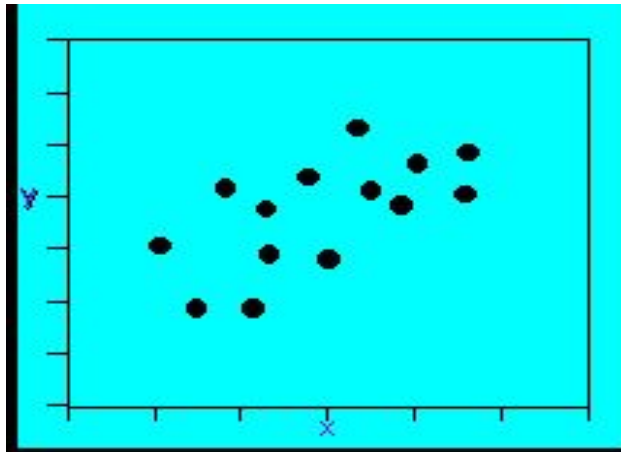


Jaya Sil

**Indian Institute of Engineering Science & Technology, Shibpur
Howrah**

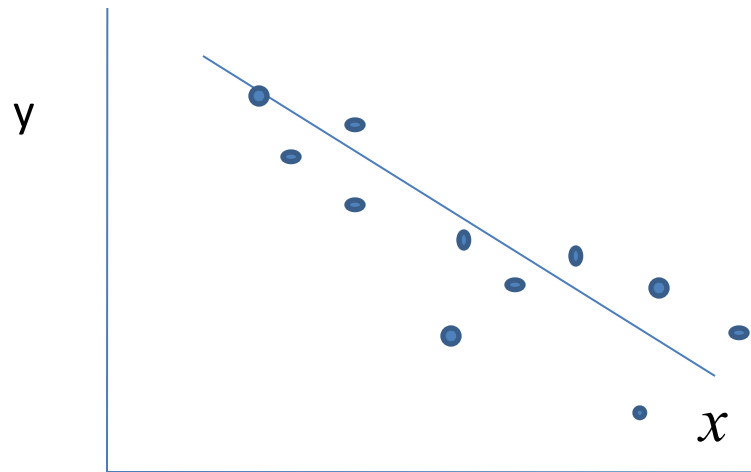
Linear Model

- **Linear models** describe a continuous response variable as a function of one or more predictor variables.
- Learning a linear relationship between the input attributes (predictor variables) and target values (response variable) values.

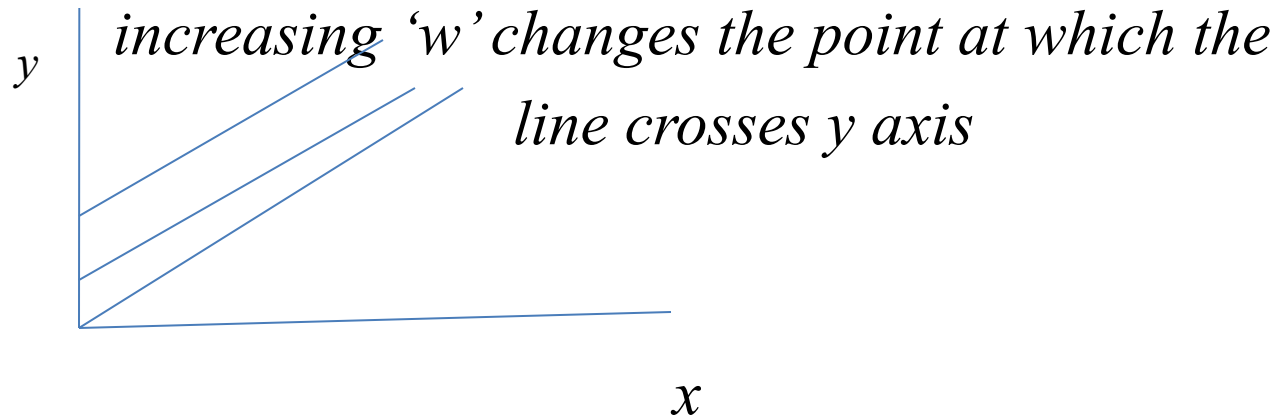


Linear Model

- Instead of evaluating $h(x)$ as a function of x , we make it more flexible using a set of associated parameters.
- $y = wx$ or $y = h(x; w)$ and the relationship between x and y is linear.
- *Assumption: The data could be adequately modeled with a straight line*



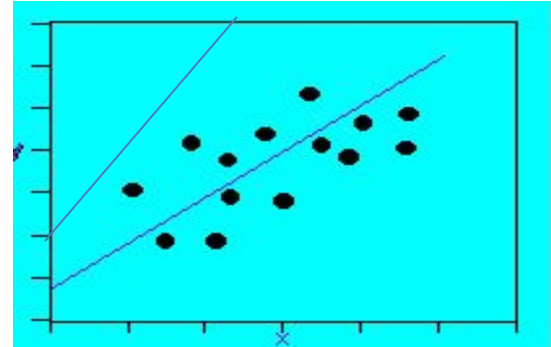
- The assumption is not perfectly satisfied in the Figure.



- *Add a single parameter as $y = wx$ or $y = h(x; w)$; enhancing the model with any gradient using the choice of w .*
- But it is not realistic at $x = 0$; $y = w \times 0$ is zero.
- Adding one more parameter to the model overcomes the problem; $y = h(x; w_0, w_1) = w_0 + w_1 x$

Supervised Machine Learning

- Increasing w_1 changes the gradient



- There are many functions which could be used to define the mapping.
- The ultimate goal is to develop a finely tuned predictor function $h(x)$ such that $y = h(x)$
- The learning task now involves using the data in figure choose two suitable values of w_0 and w_1

Supervised Machine Learning

- We decide to approximate y as a linear function of x :

$$h(x) = w_0 + w_1 x_1 + w_2 x_2 + \dots + w_n x_n$$

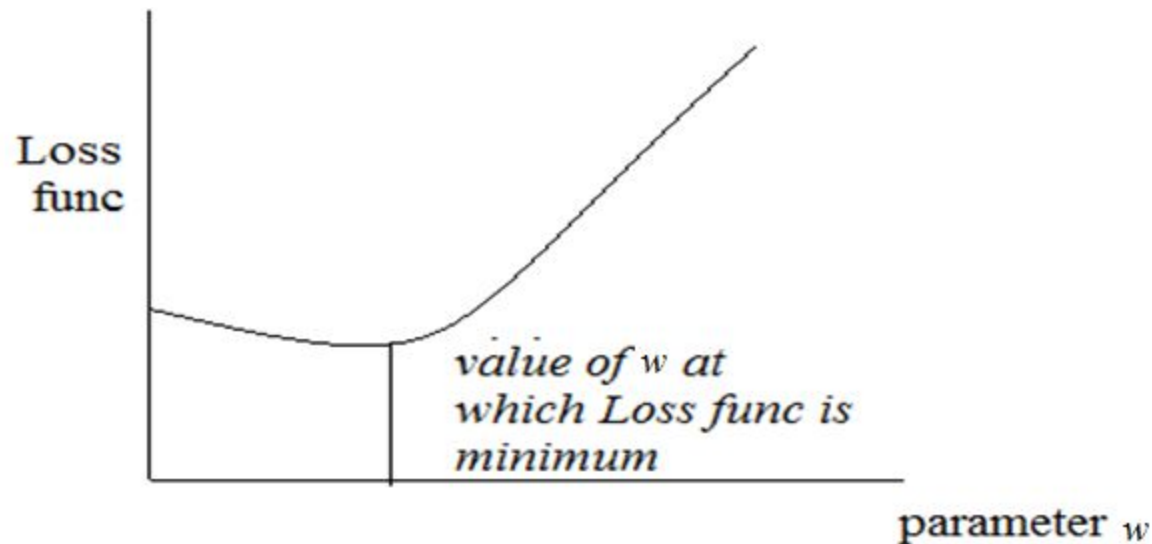
- The w_i 's are the parameters (also called weights) parameterizing the space of linear functions mapping from X to Y .
- simple predictor: $y = h(x; w_0, w_1) = w_0 + w_1 x$
Where w_0 and w_1 are constants.
- Our goal is to find the perfect values of w_0 and w_1 in order to make our predictor work as *best* as possible.
- We need to define what is the meaning of *best*.

Defining a Good Model

- The best solution consists of the values of w_0 and w_1 that produce a line that passes as close as possible to *all* of the data points.
- The minimum squared difference between the target value and the predicted value is a measure of how good is the model.
- The squared difference is defined as: $(t_n - h(x_n; w_0, w_1))^2$ for n -th pattern and known as the *squared loss function or cost function* $L_n()$
- $L_n(t_n, h(x_n; w_0, w_1)) = (t_n - h(x_n; w_0, w_1))^2$

Loss function

“Learning” optimizes the loss function so that, given input data x accurately predict value $h(x)$.



- Loss is always positive and lower the loss better the function describes the data.

- Average loss function:
$$L = \frac{1}{N} \sum_{n=1}^N L_n (t_n - h(x_n; w_0, w_1))$$

Loss Function

- Tune w_0 and w_1 to produce the model that results lowest value of the average Loss function.

- $$L = \underset{w_0, w_1}{\operatorname{arg\,min}} \quad \frac{1}{N} \sum_{n=1}^N L_n(t_n - h(x_n; w_0, w_1))$$

- Minimization of the squared loss function is the basis of Least Mean Square Error (LMSE) method of function approximation.
- Other loss functions, like Absolute Loss function

Gradient Descent Algorithm

- We want to choose w so as to minimize Loss function.
- Use a search algorithm that starts with some “initial guess” for w , and that repeatedly changes w to make Loss smaller.
- Hopefully we converge to a value of w that minimizes Loss.
- Weight updating: $w_j := w_j - \eta \frac{\partial L}{\partial w_j}$
- Weight update is simultaneously performed for all values of j . Here, η is called the learning rate.
- Gradient Descent algorithm repeatedly takes a step in the direction of steepest decrease of L .

LMS Algorithm

Tune w_0 and w_1 to produce the model that results lowest value of the average Loss function for a single training pattern.

$$\frac{\partial L}{\partial w_0} = \frac{\partial}{\partial w_0} \cdot \frac{1}{2} (t_n - h(x_n; w_0, w_1))^2$$

$$\begin{aligned} \frac{\partial L}{\partial w_0} &= (t_n - h(x_n; w_0, w_1)) \cdot \frac{\partial}{\partial w_0} (t_n - h(x_n; w_0, w_1)) \\ &= -(t_n - h(x_n; w_0, w_1)) \cdot \frac{\partial}{\partial w_0} (\sum_{n=1}^d w_n x_n - t_n) \end{aligned}$$

$$\begin{aligned} \frac{\partial L}{\partial w_1} &= (t_n - h(x_n; w_0, w_1)) \cdot \frac{\partial}{\partial w_1} (t_n - h(x_n; w_0, w_1)) \\ &= -(t_n - h(x_n; w_0, w_1)) \cdot \frac{\partial}{\partial w_1} (\sum_{n=1}^d w_n x_n - t_n) \end{aligned}$$

$$\frac{\partial L}{\partial w_n} = (h(x_n; w_0, w_1) - t_n) x_n$$

LMS UPDATE RULE

- For a single training example, the update rule is:

$$w_n := w_n - \eta (h(x_n; w_0, w_1) - t_n) x_n$$

- The magnitude of weight updating is proportional to error i.e.

$$(h(x_n; w_0, w_1) - t_n)$$

For N number of training patterns weight update rule:

$$w_n := w_n - \eta \sum_{n=1}^N (h(x_n; w_0, w_1) - t_n) x_n$$

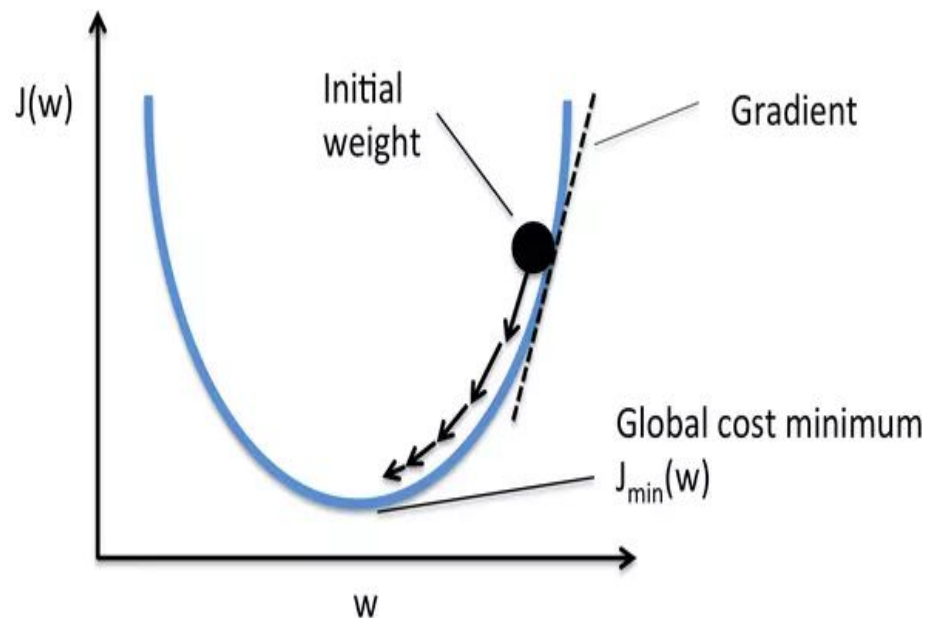
The algorithm will converge when there no weight update takes place in case it is performed iteratively.

Gradient Descent Algorithm

- The update rule is gradient descent when summation is substituted by $\frac{\partial L}{\partial w_j}$ i.e. gradient of cost or loss function.
- L is a convex quadratic function, so converges at global minima/maxima.
- When updating is performed for each training example, called Batch Gradient Descent.
- When updating is performed for a set of training example, called Stochastic Gradient descent.

- Searching for points where the gradient of a function is zero, called minima.

To determine the value of the zero gradient point (minima, maxima) we examine the second derivative



Analytical Solution

$L = 1/N \sum L_n (t_n - h(x_n; w_0, w_1))$; L is average Loss function

$$= 1/N \sum (t_n - h(x_n; w_0, w_1))^2$$

$$= 1/N \sum (t_n - (w_0 + w_1 x_n))^2$$

- Differentiating L by calculating the partial derivatives with respect to w_0 and w_1 and equating them to zero to obtain w_0 and w_1
- Differentiating again w.r.t. w_0 and w_1 we find the point at which loss is minimum.

Turning points

- $w_0 = 1/N (\sum t_n) - w_1(1/N(\sum x_n))$ when $\frac{\partial^2 L}{\partial w_0^2} = 2$
- $w_0^{av} = t^{av} - w_1 x^{av}$

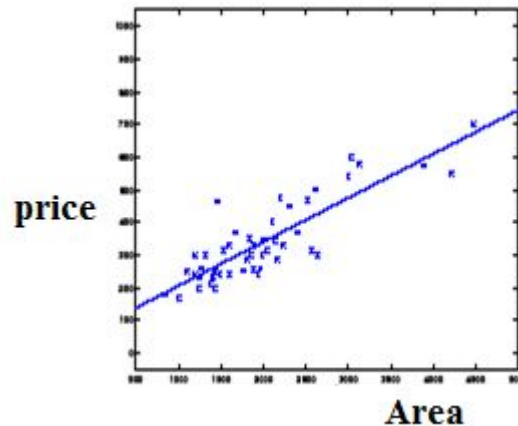
There is one turning point that correspond to minimum loss

$$w_1^{av} = \frac{\frac{1}{N} (\sum_{n=1}^N x_n t_n) - t^{av} x^{av}}{(\frac{1}{N} \sum_{n=1}^N x_n^2) - x^{av} x^{av}} \quad \text{when} \quad \frac{\partial^2 L}{\partial w_1^2} = \frac{2}{N} \sum_{n=1}^N x_n^2$$

Now we can compute the best parameter values

Prediction

- Based on linear regression model we predict the output for some input.
- A simple linear model can fit a small dataset and used for prediction.
- $w_0 = 71.27, w_1 = 0.1345$



- Linear model can be extended to larger sets of attributes, modeling complex relationship between input and output.

Vector-Matrix Notation

- Each data point is described by a set of attributes.
- Solving partial derivatives for each parameter associated with the attributes are time consuming affair.
- Representing attributes of each data point into vector form.
- For example n -th data point by \mathbf{x}_n and with two attributes
$$\mathbf{x}_n = [x_{n1}, x_{n2}]^T$$
- Column vectors \mathbf{w} and \mathbf{x}_n is defined as $h(x_n; w_0 w_1) = \mathbf{w}^T \mathbf{x}_n = w_0 + w_1 x_n$

- $L = 1/N \sum (t_n - (w_0 + w_1 x_n))^2 = 1/N \sum (t_n - \mathbf{w}^T \mathbf{x}_n)^2$
- $(\mathbf{t} - \mathbf{X}\mathbf{w})^T(\mathbf{t} - \mathbf{X}\mathbf{w})$ is used to write the loss function.
- $L = 1/N (\mathbf{t} - \mathbf{X}\mathbf{w})^T(\mathbf{t} - \mathbf{X}\mathbf{w})$
- Differentiating loss in vector/matrix form to obtain the vector \mathbf{w} corresponding to the point where L is minimum.

$$\frac{\partial L}{\partial \mathbf{w}} = \begin{bmatrix} \frac{\partial L}{\partial w_0} \\ \frac{\partial L}{\partial w_1} \end{bmatrix}$$

Making Prediction

- Given a new vector of attributes, x_{new} , the prediction using the model as $t_{\text{new}} = \mathbf{W}^T \mathbf{X}_{\text{new}}$
- Linear model of the form with multiple attributes:

$$h(x_1, x_2, \dots, x_n; w_0, w_1, \dots, w_n);$$

$$t_n = w_0 + w_1 x_{n1} + w_2 x_{n2} + \dots + ..$$

- *Prediction from such model is very precise but not always sensible.*

Learning Task

- Learning using **training examples** : statistically significant random sample.
- If the training set is too small ([law of large numbers](#)), we won't learn enough and may even reach inaccurate conclusions.
- For each training example, an input value x_{train} , and corresponding output, y or *target* is known in advance.
- For each example, we find the squared difference between the *target*, and predicted value $h(x_{\text{train}})$.
- With enough training examples, these differences give us a useful way to measure the “wrongness” of $h(x)$.

Learning Task

- Find parameter values so that the difference makes it “less wrong”.
- This process is repeated over and over until the system has converged on the best values.
- In this way, the predictor becomes trained, and is ready to do some real-world predicting.

Linear Regression

- Get familiar with objective functions, computing their gradients and optimizing the objectives over a set of parameters.
- Goal is to predict a target value y using a vector of input values $x \in \mathbb{R}^n$ where the elements x_j of x represent “features” that describe the output y .
- Suppose many examples of houses where the features for the i^{th} house are denoted $x^{(i)}$ and the price is $y^{(i)}$.
- Find a function $y = h(x)$
- If we succeed in finding a function $h(x)$ and we have seen enough examples of houses and their prices, we hope that the function $h(x)$ will also be a good predictor of the house price when we are given the features for a new house where the price is not known.

Linear Regression

$h_w(x) = \sum_j w_j x_j = w^\top x$; functions parametrized by the choice of w .

- Task is to find w so that $h_w(x^{(i)})$ is as close as possible to $y(i)$.
- In particular, we search for a w that minimizes:
$$L(w) = 1/2 \sum_i (h_w(x^{(i)}) - y^{(i)})^2 = 1/2 \sum_i (w^\top x^{(i)} - y^{(i)})^2$$
- This function is the “cost function” which measures how much error is incurred in predicting $y^{(i)}$ for a particular choice of w .
- This may also be called a “loss”, “penalty” or “objective” function.

- Find the choice of w that minimizes $L(w)$.
- The optimization procedure finds the best choice of w
- The gradient $\nabla_w L(w)$ of a differentiable function L is a vector that points in the direction of steepest increase as a function of w
- It is easy to see how an optimization algorithm could use this to make a small change to w that decreases (or increase) $L(w)$.

Optimization Method

- Compute the gradient:

- $\nabla_{\mathbf{w}} L(\mathbf{w}) =$
- $\begin{matrix} \partial L(\mathbf{w}) / \partial w_1 \\ \partial L(\mathbf{w}) / \partial w_2 \\ \vdots \\ \partial L(\mathbf{w}) / \partial w_n \end{matrix}$
-

- Differentiating the cost function $L(\mathbf{w})$ with respect to a particular parameter w_j :
- $\partial L(\mathbf{w}) / \partial w_j = \sum_i x_j^{(i)} (h_{\mathbf{w}}(x^{(i)}) - y^{(i)})$

Non-Linear Response from a Linear Model

- The linear model in terms of w and x : $h(x; w) = w_0 + w_1x$
- The model is linear in term of w only:
 $h(x; w) = w_0 + w_1x + w_2x^2$ but the function is quadratic in terms of data.
- We can add as many power we like to get a polynomial function of any order.
- The general form for a K-th order polynomial:
- $$h(x; \mathbf{w}) = \sum_{k=0}^K w_k x^k \quad OR \quad h(x; w) = w_0 + w_1x^2 + w_2x_1x_2 + w_3x_2^2 + \dots$$

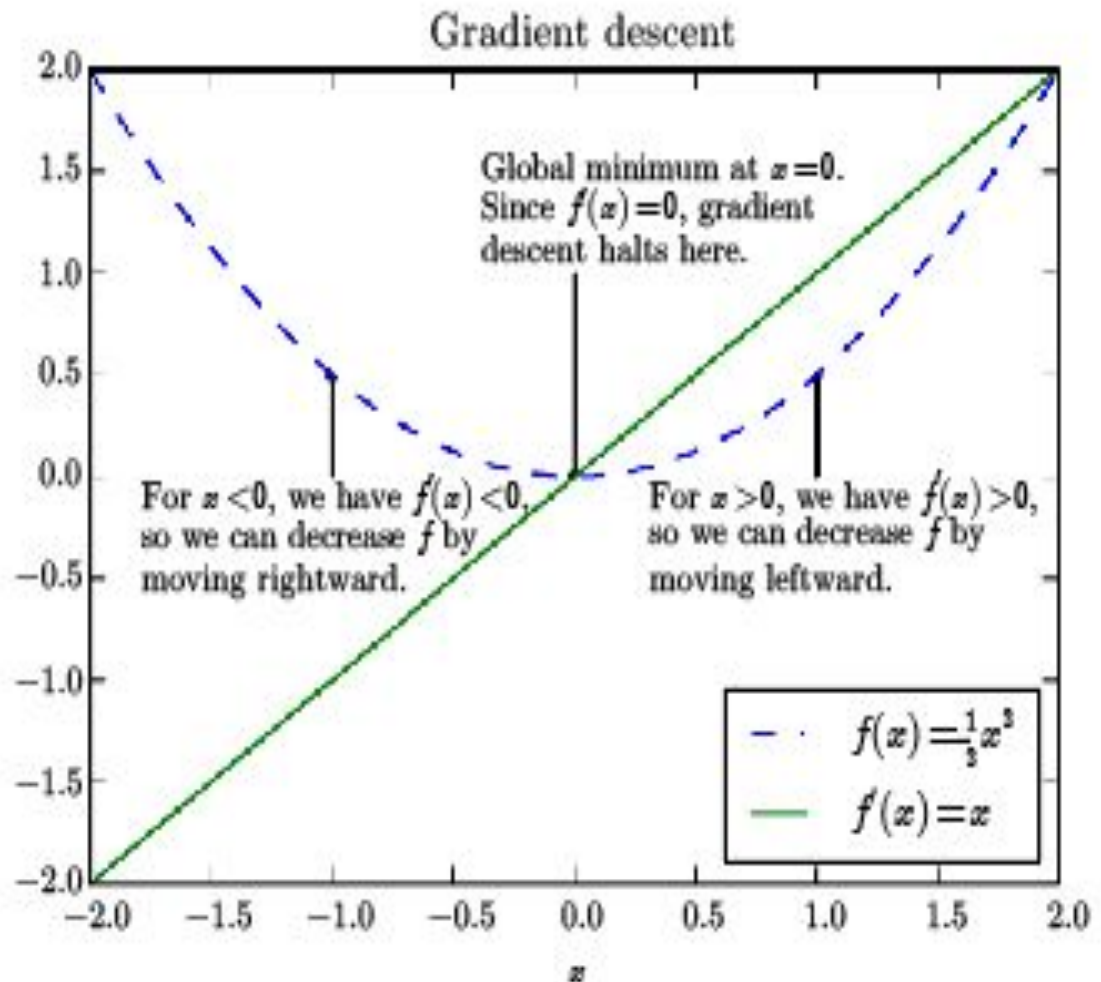
Gradient Descent

$$y = f(x)$$

x and y are
real numbers.

$$dy/dx = f'(x)$$

$f'(x)$ says how to
change x for a small
improvement of y



Critical Points

When $dy/dx = f'(x) = 0$, the derivative provides no information about which direction to move, points are called critical points.

- A local minima is a point where $y = f(x)$ is lower than at all the neighbouring points. So it is no longer possible to decrease $f(x)$ by infinitesimal steps.
- A local maxima is a point where $f(x)$ is higher than neighboring points, so not possible to increase $f(x)$

Types of critical points

